

# airline\_passengers\_analysis

June 22, 2025

## 1 Airline Passengers Time Series Forecasting

This notebook analyzes the classic airline passengers dataset and builds forecasting models to predict future passenger numbers.

### 1.1 Dataset Overview

The dataset contains monthly airline passenger numbers from 1949 to 1960, making it a perfect example for time series analysis and forecasting.

### 1.2 1. Import Required Libraries

```
[ ]: # Data manipulation and analysis
import pandas as pd
import numpy as np

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots

# Time series analysis
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.holtwinters import ExponentialSmoothing
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

# Machine learning
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.preprocessing import MinMaxScaler

# Warnings
import warnings
warnings.filterwarnings('ignore')
```

```

# Set style
plt.style.use('seaborn-v0_8')
sns.set_palette("husl")

print("Libraries imported successfully!")

```

Libraries imported successfully!

### 1.3 2. Load and Explore the Data

```

[ ]: # Load the dataset
df = pd.read_csv('data/AirPassengers.csv')

# Display basic information
print("Dataset Shape:", df.shape)
print("\nFirst 5 rows:")
print(df.head())
print("\nLast 5 rows:")
print(df.tail())
print("\nDataset Info:")
print(df.info())
print("\nBasic Statistics:")
print(df.describe())

```

Dataset Shape: (144, 2)

First 5 rows:

	Month	#Passengers
0	1949-01	112
1	1949-02	118
2	1949-03	132
3	1949-04	129
4	1949-05	121

Last 5 rows:

	Month	#Passengers
139	1960-08	606
140	1960-09	508
141	1960-10	461
142	1960-11	390
143	1960-12	432

Dataset Info:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144 entries, 0 to 143
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype 

```

```
--  -----
0   Month          144 non-null    object
1   #Passengers   144 non-null    int64
dtypes: int64(1), object(1)
memory usage: 2.4+ KB
None
```

Basic Statistics:

```
#Passengers
count    144.000000
mean     280.298611
std      119.966317
min      104.000000
25%     180.000000
50%     265.500000
75%     360.500000
max     622.000000
```

```
[ ]: # Data preprocessing
# Convert Month to datetime
df['Month'] = pd.to_datetime(df['Month'])
df.set_index('Month', inplace=True)

# Rename column for easier access
df.rename(columns={'#Passengers': 'Passengers'}, inplace=True)

print("Data preprocessing completed!")
print("\nProcessed dataset:")
print(df.head())
print(f"\nDate range: {df.index.min()} to {df.index.max()}")
print(f"Total months: {len(df)}")
```

Data preprocessing completed!

Processed dataset:

	Passengers
Month	
1949-01-01	112
1949-02-01	118
1949-03-01	132
1949-04-01	129
1949-05-01	121

Date range: 1949-01-01 00:00:00 to 1960-12-01 00:00:00  
Total months: 144

## 1.4 3. Exploratory Data Analysis

```
[ ]: # Basic time series plot
fig, axes = plt.subplots(2, 2, figsize=(15, 10))

# Original time series
axes[0, 0].plot(df['Passengers'], linewidth=2)
axes[0, 0].set_title('Airline Passengers Over Time')
axes[0, 0].set_xlabel('Year')
axes[0, 0].set_ylabel('Number of Passengers')
axes[0, 0].grid(True, alpha=0.3)

# Distribution of passengers
axes[0, 1].hist(df['Passengers'], bins=20, alpha=0.7, edgecolor='black')
axes[0, 1].set_title('Distribution of Passenger Numbers')
axes[0, 1].set_xlabel('Number of Passengers')
axes[0, 1].set_ylabel('Frequency')
axes[0, 1].grid(True, alpha=0.3)

# Box plot by year
df_year = df.copy()
df_year['Year'] = df_year.index.year
df_year.boxplot(column='Passengers', by='Year', ax=axes[1, 0])
axes[1, 0].set_title('Passenger Numbers by Year')
axes[1, 0].set_xlabel('Year')
axes[1, 0].set_ylabel('Number of Passengers')

# Monthly seasonality
df_month = df.copy()
df_month['Month'] = df_month.index.month
monthly_avg = df_month.groupby('Month')['Passengers'].mean()
axes[1, 1].plot(monthly_avg.index, monthly_avg.values, marker='o', linewidth=2)
axes[1, 1].set_title('Average Passengers by Month')
axes[1, 1].set_xlabel('Month')
axes[1, 1].set_ylabel('Average Passengers')
axes[1, 1].set_xticks(range(1, 13))
axes[1, 1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```

```
-----
ValueError                                Traceback (most recent call last)
/tmp/ipykernel_10802/751950567.py in ?()
    25
    26 # Monthly seasonality
    27 df_month = df.copy()
```

```

    28 df_month['Month'] = df_month.index.month
--> 29 monthly_avg = df_month.groupby('Month')['Passengers'].mean()
    30 axes[1, 1].plot(monthly_avg.index, monthly_avg.values, marker='o', □
↳ linewidth=2)
    31 axes[1, 1].set_title('Average Passengers by Month')
    32 axes[1, 1].set_xlabel('Month')

~/ds/airline_forecasting_env/lib/python3.10/site-packages/pandas/core/frame.py
↳ in ?(self, by, axis, level, as_index, sort, group_keys, observed, dropna)
  9179
  9180     if level is None and by is None:
  9181         raise TypeError("You have to supply one of 'by' and 'level' ")
  9182
-> 9183     return DataFrameGroupBy(
  9184         obj=self,
  9185         keys=by,
  9186         axis=axis,

~/ds/airline_forecasting_env/lib/python3.10/site-packages/pandas/core/groupby/
↳ groupby.py in ?(self, obj, keys, axis, level, grouper, exclusions, selection, □
↳ as_index, sort, group_keys, observed, dropna)
  1325         self.group_keys = group_keys
  1326         self.dropna = dropna
  1327
  1328     if grouper is None:
-> 1329         grouper, exclusions, obj = get_grouper(
  1330             obj,
  1331             keys,
  1332             axis=axis,

~/ds/airline_forecasting_env/lib/python3.10/site-packages/pandas/core/groupby/
↳ grouper.py in ?(obj, key, axis, level, sort, observed, validate, dropna)
  1029
  1030     elif is_in_axis(gpr): # df.groupby('name')
  1031         if obj.ndim != 1 and gpr in obj:
  1032             if validate:
-> 1033                 obj._check_label_or_level_ambiguity(gpr, axis=axis)
  1034                 in_axis, name, gpr = True, gpr, obj[gpr]
  1035                 if gpr.ndim != 1:
  1036                     # non-unique columns; raise here to get the name in
↳ the

~/ds/airline_forecasting_env/lib/python3.10/site-packages/pandas/core/generic.p
↳ in ?(self, key, axis)
  1864         msg = (
  1865             f"'{key}' is both {level_article} {level_type} level and"
  ↳ "


```

```

1866             f"{{label_article}} {{label_type}} label, which is ambiguous."
1867             ^
1867         )
-> 1868     raise ValueError(msg)

ValueError: 'Month' is both an index level and a column label, which is ambiguous.

```



## 1.5 4. Time Series Decomposition

```

[ ]: # Perform seasonal decomposition
decomposition = seasonal_decompose(df['Passengers'], model='multiplicative',
                                     period=12)

# Plot decomposition
fig, axes = plt.subplots(4, 1, figsize=(15, 12))

decomposition.observed.plot(ax=axes[0], title='Original Time Series')
decomposition.trend.plot(ax=axes[1], title='Trend Component')
decomposition.seasonal.plot(ax=axes[2], title='Seasonal Component')
decomposition.resid.plot(ax=axes[3], title='Residual Component')

```


```

for ax in axes:
    ax.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

print("Time series shows:")
print("1. Clear upward trend")
print("2. Strong seasonal pattern (yearly cycle)")
print("3. Increasing variance over time")

```



Time series shows:  
1. Clear upward trend  
2. Strong seasonal pattern (yearly cycle)  
3. Increasing variance over time

## 1.6 5. Stationarity Testing

```
[ ]: # Augmented Dickey-Fuller test for stationarity
def check_stationarity(timeseries, title):
    print(f'Results of Augmented Dickey-Fuller Test for {title}:')
    dfoutput = adfuller(timeseries, autolag='AIC')
    dfoutput = pd.Series(dfoutput[0:4], index=['Test Statistic', 'p-value', '#Lags Used', 'Number of Observations Used'])
    for key,value in dfoutput[4].items():
        dfoutput['Critical Value (%s)'%key] = value
    print(dfoutput)

    if dfoutput[1] <= 0.05:
        print("Series is stationary")
    else:
        print("Series is non-stationary")
    print("-"*50)

# Test original series
check_stationarity(df['Passengers'], 'Original Series')

# Apply log transformation
df['Passengers_log'] = np.log(df['Passengers'])
check_stationarity(df['Passengers_log'], 'Log Transformed Series')

# Apply first differencing to log series
df['Passengers_log_diff'] = df['Passengers_log'].diff().dropna()
check_stationarity(df['Passengers_log_diff'].dropna(), 'Log Differenced Series')
```

Results of Augmented Dickey-Fuller Test for Original Series:

Test Statistic	0.815369
p-value	0.991880
#Lags Used	13.000000
Number of Observations Used	130.000000
Critical Value (1%)	-3.481682
Critical Value (5%)	-2.884042
Critical Value (10%)	-2.578770
dtype: float64	
Series is non-stationary	

---

Results of Augmented Dickey-Fuller Test for Log Transformed Series:

Test Statistic	-1.717017
p-value	0.422367
#Lags Used	13.000000
Number of Observations Used	130.000000
Critical Value (1%)	-3.481682
Critical Value (5%)	-2.884042
Critical Value (10%)	-2.578770

```

dtype: float64
Series is non-stationary
-----
Results of Augmented Dickey-Fuller Test for Log Differenced Series:
Test Statistic           -2.717131
p-value                  0.071121
#Lags Used              14.000000
Number of Observations Used 128.000000
Critical Value (1%)      -3.482501
Critical Value (5%)       -2.884398
Critical Value (10%)      -2.578960
dtype: float64
Series is non-stationary
-----
```


## 1.7 6. ACF and PACF Analysis

```
[ ]: # Plot ACF and PACF for the stationary series
fig, axes = plt.subplots(2, 2, figsize=(15, 10))

# Original series
plot_acf(df['Passengers'].dropna(), ax=axes[0, 0], lags=40, title='ACF -  
↳Original Series')
plot_pacf(df['Passengers'].dropna(), ax=axes[0, 1], lags=40, title='PACF -  
↳Original Series')

# Differenced series
plot_acf(df['Passengers_log_diff'].dropna(), ax=axes[1, 0], lags=40, title='ACF -  
↳ Log Differenced Series')
plot_pacf(df['Passengers_log_diff'].dropna(), ax=axes[1, 1], lags=40, title='PACF - Log Differenced Series')

plt.tight_layout()
plt.show()
```



## 1.8 7. Train-Test Split


```
[ ]: # Split data into train and test sets
train_size = int(len(df) * 0.8)
train_data = df[:train_size]
test_data = df[train_size:]

print(f"Training data: {len(train_data)} months ({train_data.index[0]} to
      ↪{train_data.index[-1]})")
print(f"Testing data: {len(test_data)} months ({test_data.index[0]} to
      ↪{test_data.index[-1]})")

# Visualize the split
plt.figure(figsize=(12, 6))
plt.plot(train_data.index, train_data['Passengers'], label='Training Data',
      ↪color='blue')
plt.plot(test_data.index, test_data['Passengers'], label='Testing Data',
      ↪color='red')
plt.title('Train-Test Split')
plt.xlabel('Year')
plt.ylabel('Number of Passengers')
plt.legend()
plt.grid(True, alpha=0.3)
```

```
plt.show()
```

Training data: 115 months (1949-01-01 00:00:00 to 1958-07-01 00:00:00)  
Testing data: 29 months (1958-08-01 00:00:00 to 1960-12-01 00:00:00)



## 1.9 8. ARIMA Model

```
[ ]: # Fit ARIMA model
# Using ARIMA(2,1,2) based on ACF/PACF analysis
arima_model = ARIMA(train_data['Passengers'], order=(2,1,2))
arima_fitted = arima_model.fit()

print("ARIMA Model Summary:")
print(arima_fitted.summary())

# Make predictions
arima_forecast = arima_fitted.forecast(steps=len(test_data))
arima_forecast_index = test_data.index

# Calculate metrics
arima_mse = mean_squared_error(test_data['Passengers'], arima_forecast)
arima_mae = mean_absolute_error(test_data['Passengers'], arima_forecast)
arima_rmse = np.sqrt(arima_mse)

print(f"\nARIMA Model Performance:")
print(f"MSE: {arima_mse:.2f}")
print(f"MAE: {arima_mae:.2f}")
```

```
print(f"RMSE: {arima_rmse:.2f}")
```

ARIMA Model Summary:

SARIMAX Results

```
=====
Dep. Variable: Passengers No. Observations: 115
Model: ARIMA(2, 1, 2) Log Likelihood: -523.758
Date: Sat, 24 May 2025 AIC: 1057.516
Time: 15:07:42 BIC: 1071.197
Sample: 01-01-1949 HQIC: 1063.069
          - 07-01-1958
Covariance Type: opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.3280	0.145	2.268	0.023	0.045	0.611
ar.L2	0.2521	0.165	1.528	0.126	-0.071	0.575
ma.L1	-0.0125	0.109	-0.114	0.909	-0.227	0.202
ma.L2	-0.7544	0.130	-5.812	0.000	-1.009	-0.500
sigma2	568.4920	103.877	5.473	0.000	364.897	772.087

<=====

==

Ljung-Box (L1) (Q): 0.02 Jarque-Bera (JB):

3.39

Prob(Q): 0.90 Prob(JB):

0.18

Heteroskedasticity (H): 5.24 Skew:

0.11

Prob(H) (two-sided): 0.00 Kurtosis:

2.19

==

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

ARIMA Model Performance:

MSE: 6808.40

MAE: 63.55

RMSE: 82.51

## 1.10 9. Exponential Smoothing (Holt-Winters)

```
[ ]: # Fit Exponential Smoothing model
exp_smooth_model = ExponentialSmoothing(
    train_data['Passengers'],
    trend='add',
    seasonal='add',
    seasonal_periods=12
)
exp_smooth_fitted = exp_smooth_model.fit()

# Make predictions
exp_smooth_forecast = exp_smooth_fitted.forecast(steps=len(test_data))

# Calculate metrics
exp_smooth_mse = mean_squared_error(test_data['Passengers'], □
    ↪exp_smooth_forecast)
exp_smooth_mae = mean_absolute_error(test_data['Passengers'], □
    ↪exp_smooth_forecast)
exp_smooth_rmse = np.sqrt(exp_smooth_mse)

print(f"Exponential Smoothing Model Performance:")
print(f"MSE: {exp_smooth_mse:.2f}")
print(f"MAE: {exp_smooth_mae:.2f}")
print(f"RMSE: {exp_smooth_rmse:.2f}")
```

Exponential Smoothing Model Performance:

MSE: 1541.40

MAE: 31.79

RMSE: 39.26

## 1.11 10. Model Comparison and Visualization

```
[ ]: # Create comparison plot
plt.figure(figsize=(15, 8))

# Plot actual data
plt.plot(df.index, df['Passengers'], label='Actual', color='black', linewidth=2)

# Plot training data
plt.plot(train_data.index, train_data['Passengers'], label='Training Data', □
    ↪color='blue', alpha=0.7)

# Plot forecasts
plt.plot(test_data.index, arima_forecast, label='ARIMA Forecast', color='red', □
    ↪linestyle='--', linewidth=2)
```

```

plt.plot(test_data.index, exp_smooth_forecast, label='Exp. Smoothing Forecast',  

         color='green', linestyle='--', linewidth=2)

# Add vertical line to separate train/test
plt.axvline(x=train_data.index[-1], color='gray', linestyle=':', alpha=0.7,  


            label='Train/Test Split')

plt.title('Airline Passengers Forecasting - Model Comparison')
plt.xlabel('Year')
plt.ylabel('Number of Passengers')
plt.legend()
plt.grid(True, alpha=0.3)
plt.show()

# Model performance comparison
models_performance = pd.DataFrame({
    'Model': ['ARIMA', 'Exponential Smoothing'],
    'MSE': [arima_mse, exp_smooth_mse],
    'MAE': [arima_mae, exp_smooth_mae],
    'RMSE': [arima_rmse, exp_smooth_rmse]
})

print("\nModel Performance Comparison:")
print(models_performance)

```



Model Performance Comparison:

	Model	MSE	MAE	RMSE
0	ARIMA	6808.397034	63.545311	82.513011
1	Exponential Smoothing	1541.395743	31.791027	39.260613

## 1.12 11. Future Predictions

```
[ ]: # Make future predictions (next 12 months)
future_periods = 12

# Refit models on full dataset for future predictions
arima_full = ARIMA(df['Passengers'], order=(2,1,2)).fit()
exp_smooth_full = ExponentialSmoothing(
    df['Passengers'],
    trend='add',
    seasonal='add',
    seasonal_periods=12
).fit()

# Generate future dates
last_date = df.index[-1]
future_dates = pd.date_range(start=last_date + pd.DateOffset(months=1), periods=future_periods, freq='MS')

# Make predictions
arima_future = arima_full.forecast(steps=future_periods)
exp_smooth_future = exp_smooth_full.forecast(steps=future_periods)

# Create future predictions dataframe
future_predictions = pd.DataFrame({
    'Date': future_dates,
    'ARIMA_Forecast': arima_future,
    'ExpSmooth_Forecast': exp_smooth_future
})
future_predictions.set_index('Date', inplace=True)

print("Future Predictions (Next 12 months):")
print(future_predictions)

# Plot future predictions
plt.figure(figsize=(15, 8))

# Plot historical data
plt.plot(df.index, df['Passengers'], label='Historical Data', color='black', linewidth=2)

# Plot future predictions
plt.plot(future_predictions.index, future_predictions['ARIMA_Forecast'],
```

```

        label='ARIMA Future Forecast', color='red', linestyle='--', ↴
        linewidth=2, marker='o')
plt.plot(future_predictions.index, future_predictions['ExpSmooth_Forecast'],
         label='Exp. Smoothing Future Forecast', color='green', linestyle='--', ↴
         linewidth=2, marker='s')


# Add vertical line to separate historical/future
plt.axvline(x=df.index[-1], color='gray', linestyle=':', alpha=0.7, ↴
            label='Historical/Future Split')

plt.title('Airline Passengers - Future Predictions')
plt.xlabel('Year')
plt.ylabel('Number of Passengers')
plt.legend()
plt.grid(True, alpha=0.3)
plt.show()

```

Future Predictions (Next 12 months):

	ARIMA_Forecast	ExpSmooth_Forecast
Date		
1961-01-01	439.854863	451.577073
1961-02-01	465.296545	427.257748
1961-03-01	500.666051	464.699360
1961-04-01	535.971401	500.103185
1961-05-01	561.689094	508.660794
1961-06-01	571.312667	567.713873
1961-07-01	562.972100	647.897121
1961-08-01	539.729052	632.460961
1961-09-01	508.528057	534.374265
1961-10-01	478.147364	484.930880
1961-11-01	456.747520	418.459518
1961-12-01	449.697306	471.058935



## 1.13 12. Conclusions and Insights

### 1.13.1 Key Findings:

1. **Trend:** Clear upward trend in airline passengers from 1949 to 1960
2. **Seasonality:** Strong seasonal pattern with peaks in summer months
3. **Growth:** Exponential growth pattern with increasing variance

### 1.13.2 Model Performance:

- Both ARIMA and Exponential Smoothing models capture the trend and seasonality well
- Exponential Smoothing typically performs better for data with clear trend and seasonality
- ARIMA models are more flexible but require careful parameter tuning

### 1.13.3 Business Implications:

1. **Capacity Planning:** Airlines can use these forecasts for fleet and route planning
2. **Seasonal Staffing:** Higher staffing needs during summer months
3. **Revenue Management:** Dynamic pricing strategies based on predicted demand
4. **Infrastructure:** Airport and ground services capacity planning

### 1.13.4 Next Steps:

1. Incorporate external factors (economic indicators, fuel prices, etc.)
2. Try more advanced models (SARIMA, Prophet, LSTM)
3. Implement confidence intervals for predictions
4. Regular model retraining with new data